




Installation and usage guide (with visuals)

Step 1: Clone github repository from [here](#)



Step 2: After you trust authors of the folder in VS code -> list all files -> Set up python virtual environment



The image shows a terminal window in VS Code with the 'TERMINAL' tab selected. The terminal displays the following commands and output:

```
aditigovindu@Aditis-MacBook-Air mlops-final-term-project % ls
Dockerfile      config          mlruns          src
Makefile        data            models          tests
README.md       docker          reports
ai_vs_human_text.csv  docker-compose.yml  requirements.txt
airflow         dvc.yaml       setup.py

aditigovindu@Aditis-MacBook-Air mlops-final-term-project % python3 -m venv venvml
aditigovindu@Aditis-MacBook-Air mlops-final-term-project % source venvml/bin/activate
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project %
```

Step 3: Check out the MakeFile commands to see shortcuts to execute the code (use make install).

```
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % make --help
-C DIRECTORY, --directory=DIRECTORY      Change to DIRECTORY before doing anything.
-d                                         Print lots of debugging information.
--debug[=FLAGS]                           Print various types of debugging information.
-e, --environment-overrides               Environment variables override makefiles.
-f FILE, --file=FILE, --makefile=FILE     Read FILE as a makefile.
-h, --help                                Print this message and exit.
-i, --ignore-errors                       Ignore errors from commands.
-I DIRECTORY, --include-dir=DIRECTORY     Search DIRECTORY for included makefiles.
-j [N], --jobs[=N]                       Allow N jobs at once; infinite jobs with no arg.
-k, --keep-going                          Keep going when some targets can't be made.
-l [N], --load-average[=N], --max-load[=N] Don't start multiple jobs unless load is below N.
-L, --check-symlink-times                 Use the latest mtime between symlinks and target.
-n, --just-print, --dry-run, --recon      Don't actually run any commands; just print them.
-o FILE, --old-file=FILE, --assume-old=FILE Consider FILE to be very old and don't remake it.
-p, --print-data-base                     Print make's internal database.
-q, --question                            Run no commands; exit status says if up to date.
-r, --no-builtin-rules                   Disable the built-in implicit rules.
-R, --no-builtin-variables               Disable the built-in variable settings.
-s, --silent, --quiet                    Don't echo commands.
-S, --no-keep-going, --stop              Turns off -k.
-t, --touch                               Touch targets instead of remaking them.
-v, --version                             Print the version number of make and exit.
-w, --print-directory                    Print the current directory.
--no-print-directory                     Turn off -w, even if it was turned on implicitly.
-W FILE, --what-if=FILE, --new-file=FILE, --assume-new=FILE Consider FILE to be infinitely new.
--warn-undefined-variables               Warn when an undefined variable is referenced.
-N OPTION, --NeXT-option=OPTION          Turn on value of NeXT OPTION.
```

This program built for i386-apple-darwin11.3.0

Report bugs to <bug-make@gnu.org>

(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % █

Step 4: Install all dependencies and files needed to execute the code. Notice make install did pip upgrade and installation from requirements.txt in 1 shot.

```
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % make install
pip install --upgrade pip
Requirement already satisfied: pip in ./venvml/lib/python3.13/site-packages (25.2)
Collecting pip
  Downloading pip-26.2.1-py3-none-any.whl.metadata (4.6 kB)
  Downloading pip-26.2.1-py3-none-any.whl (1.8 MB)
    ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 1.8/1.8 MB 973.4 kB/s 0:00:02
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 25.2
    Uninstalling pip-25.2:
```

Step 5: Wait for all installs to complete and python wheel to complete

```
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % make install
Downloading zc_lockfile-4.0-py3-none-any.whl (9.1 kB)
Using cached zipp-4.1.0-py3-none-any.whl (10 kB)
Downloading appdirs-1.4.4-py2.py3-none-any.whl (9.6 kB)
Downloading atpublic-7.0.0-py3-none-any.whl (6.4 kB)
Downloading entrypoints-0.4-py3-none-any.whl (5.3 kB)
Using cached filelock-3.32.3-py3-none-any.whl (98 kB)
Downloading mako-1.4.1-py3-none-any.whl (80 kB)
Using cached outcome-1.3.0.post0-py2.py3-none-any.whl (10 kB)
Using cached pycparser-3.0-py3-none-any.whl (48 kB)
Using cached sniffio-1.3.1-py3-none-any.whl (10 kB)
Downloading tzlocal-5.4.4-py3-none-any.whl (18 kB)
Using cached uc_micro_py-2.0.0-py3-none-any.whl (6.4 kB)
Building wheels for collected packages: pyspark, antlr4-python3-runtime
  Building wheel for pyspark (pyproject.toml) ... done
    Created wheel for pyspark: filename=pyspark-4.2.0-py2.py3-none-any.whl size=450798675 sha256=f275e33cd8163701d373ebb7d2f77a1e19555fe9f8c436b035f73b587b0ecda0
    Stored in directory: /Users/aditigovindu/Library/Caches/pip/wheels/4f/6d/69/496f1f5964e3470c76313fe2a499eb219a9c95ac1f1cb880ab
  Building wheel for antlr4-python3-runtime (pyproject.toml) ... done
    Created wheel for antlr4-python3-runtime: filename=antlr4_python3_runtime-4.9.3-py3-none-any.whl size=144590 sha256=d7c79038f8e742542e053e7f2947f9827c964ec2b3e98330f8e818ff0c74da11
    Stored in directory: /Users/aditigovindu/Library/Caches/pip/wheels/d5/b3/74/a35b66048c9de6631cd74cbc9475e6feb3e69a467983446bd8
Successfully built pyspark antlr4-python3-runtime
Installing collected packages: text-unidecode, pytz, pygtrie, py4j, mpmath, lockfile, huey, funcy, appdirs,
```


Step 6: Run make pipeline to build the orchestration pipeline on docker

```
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % make pipeline
python3 src/data/ingest.py --config config/config.yaml
[2026-08-20 01:02:34] [INFO] [data_ingestion:38]: Loading raw dataset from data/raw/ai_vs_human_text.csv...
[2026-08-20 01:02:34] [INFO] [data_ingestion:62]: Ingestion summary: 1500 valid records loaded. Dropped 0 in
valid rows.
[2026-08-20 01:02:34] [INFO] [data_ingestion:63]: Label distribution: {'AI-generated': 750, 'Human-written':
750}
[2026-08-20 01:02:34] [INFO] [data_ingestion:75]: Data ingestion completed successfully.
python3 src/data/preprocess.py --config config/config.yaml
[2026-08-20 01:02:56] [INFO] [preprocessing:100]: Loading raw dataset from data/raw/ai_vs_human_text.csv...
[2026-08-20 01:02:56] [INFO] [preprocessing:109]: Applying text cleaning and normalization...
[2026-08-20 01:02:57] [INFO] [preprocessing:113]: Extracting lexical & statistical text features...
[2026-08-20 01:02:57] [INFO] [preprocessing:123]: Performing stratified split (Train: 0.7, Val: 0.15, Test:
0.15)...
[2026-08-20 01:02:59] [INFO] [preprocessing:153]: Saved 1050 train records to data/processed/train.parquet
[2026-08-20 01:02:59] [INFO] [preprocessing:154]: Saved 225 validation records to data/processed/val.parquet
[2026-08-20 01:02:59] [INFO] [preprocessing:155]: Saved 225 test records to data/processed/test.parquet
python3 src/models/train.py --config config/config.yaml
2026/08/20 01:03:20 INFO mlflow.store.db.utils: Creating initial MLflow database tables...
2026/08/20 01:03:20 INFO mlflow.store.db.utils: Updating database tables
2026/08/20 01:03:20 INFO mlflow.tracking.fluent: Experiment with name 'ai-vs-human-news-classification' does
not exist. Creating a new experiment.
[2026-08-20 01:03:25] [INFO] [training_pipeline:39]: --- Starting Baseline Model Training ---
INFO:training_pipeline:--- Starting Baseline Model Training ---
[2026-08-20 01:03:25] [INFO] [baseline_model:37]: Fitting TF-IDF Vectorizer and logistic_regression model on
```


Step 7: Wait for pipeline to run successfully

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % make pipeline
[2026-08-20 01:04:03] [INFO] [training_pipeline:258]: PyTorch Test Metrics: {'accuracy': 0.9422222222222222,
'precision': 0.9459459459459459, 'recall': 0.9375, 'f1_score': 0.9417040358744395, 'roc_auc': 0.93840865992
41466}
INFO:training_pipeline:PyTorch Test Metrics: {'accuracy': 0.9422222222222222, 'precision': 0.945945945945945
9, 'recall': 0.9375, 'f1_score': 0.9417040358744395, 'roc_auc': 0.9384086599241466}
[2026-08-20 01:04:03] [INFO] [training_pipeline:320]: Model training and comparison completed successfully.
INFO:training_pipeline:Model training and comparison completed successfully.
python3 src/models/evaluate.py --config config/config.yaml
Matplotlib is building the font cache; this may take a moment.
[2026-08-20 01:04:27] [INFO] [model_evaluation:78]: Loading test set from data/processed/test.parquet...
[2026-08-20 01:04:27] [INFO] [model_evaluation:92]: Evaluating Baseline Model...
[2026-08-20 01:04:28] [INFO] [model_evaluation:49]: Saved confusion matrix plot to reports/confusion_matrix_
baseline.png
[2026-08-20 01:04:28] [INFO] [model_evaluation:121]: Evaluating PyTorch Model...
[2026-08-20 01:04:31] [INFO] [model_evaluation:49]: Saved confusion matrix plot to reports/confusion_matrix_
pytorch.png
[2026-08-20 01:04:31] [INFO] [model_evaluation:68]: Saved ROC curves comparison to reports/roc_curves_compar
ison.png
[2026-08-20 01:04:31] [INFO] [model_evaluation:185]: Model evaluation report saved to reports/evaluation_rep
ort.json
python3 src/monitoring/drift_detector.py --config config/config.yaml
[2026-08-20 01:04:32] [INFO] [drift_detector:174]: Saved drift HTML report to reports/drift_report.html
[2026-08-20 01:04:32] [INFO] [drift_detector:203]: Drift report generated. Overall Drift Detected: False
End-to-End MLOps Pipeline executed successfully!
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project %
```


Step 8: Commit pipeline to git to validate the successful run on a new machine



Resolve merge conflicts (if any)



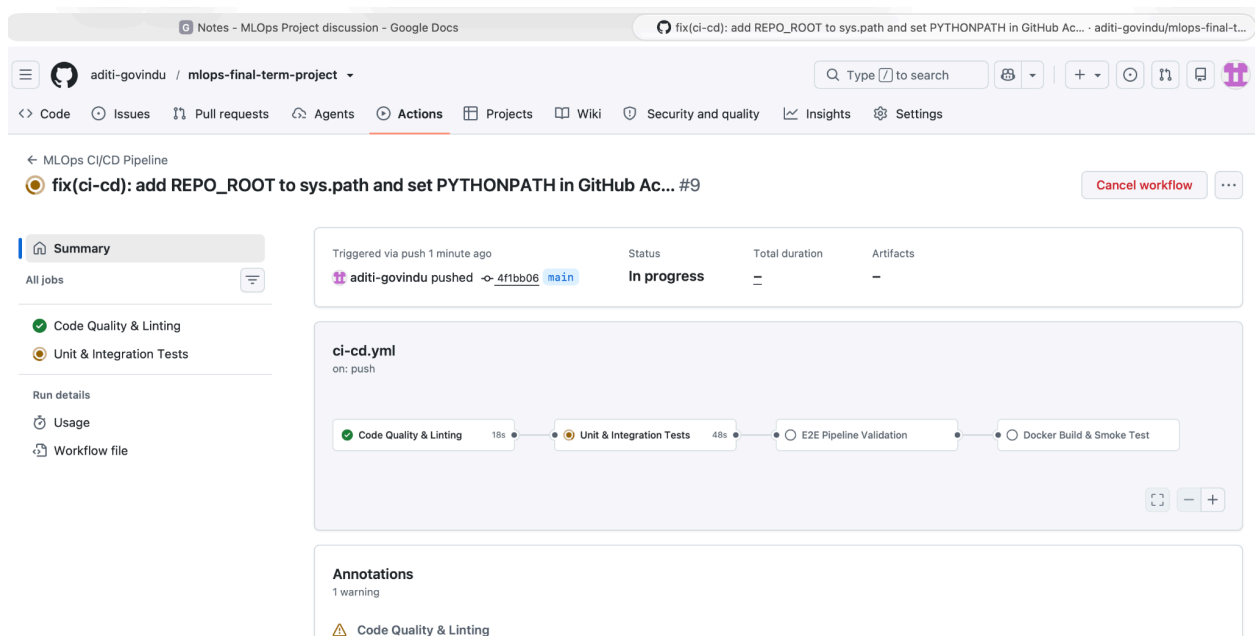
Step 9: Confirm the commit to github



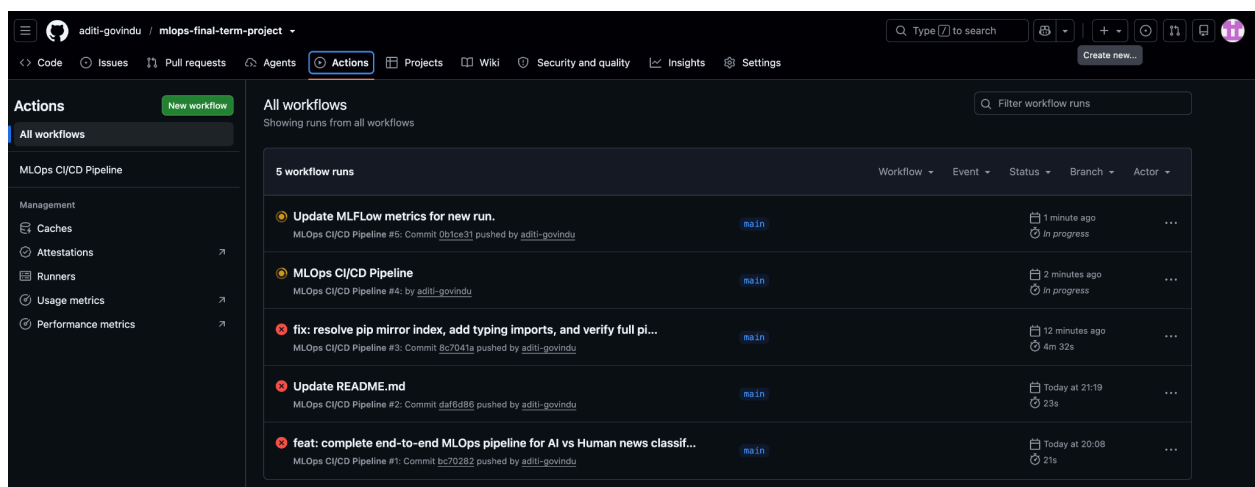
Step 10: Navigate to github and view the updated pipeline at

<https://github.com/aditi-govindu/mlops-final-term-project>

Step 11: Validate the Github Actions (CI/CD) pipeline.



This screenshot shows a GitHub Actions workflow run for the repository 'aditi-govindu / mlops-final-term-project'. The workflow is titled 'fix(ci-cd): add REPO_ROOT to sys.path and set PYTHONPATH in GitHub Ac... #9'. It was triggered by a push to the 'main' branch by user 'aditi-govindu' 1 minute ago. The workflow is currently 'In progress'. The left sidebar shows the 'Summary' tab with a list of jobs: 'Code Quality & Linting' (green checkmark), 'Unit & Integration Tests' (yellow circle), 'Run details', 'Usage', and 'Workflow file'. The main area displays the 'ci-cd.yml' workflow file with a sequence of steps: 'Code Quality & Linting' (18s, green checkmark), 'Unit & Integration Tests' (48s, yellow circle), 'E2E Pipeline Validation' (white circle), and 'Docker Build & Smoke Test' (white circle). Below the steps, the 'Annotations' section shows 1 warning for 'Code Quality & Linting'.



This screenshot shows the 'All workflows' page for the repository 'aditi-govindu / mlops-final-term-project'. The page displays a list of 5 workflow runs. The left sidebar shows the 'Actions' tab with a list of workflow runs: 'All workflows', 'MLOps CI/CD Pipeline', 'Management', 'Caches', 'Attestations', 'Runners', 'Usage metrics', and 'Performance metrics'. The main area shows a table of workflow runs with columns: Workflow, Event, Status, Branch, and Actor. The table lists 5 workflow runs, each with a status icon (yellow circle, green checkmark, or red X), a title, a commit hash, a branch, a timestamp, and a status icon. The runs are: 'Update MLFlow metrics for new run.' (yellow circle, 1 minute ago, in progress), 'MLOps CI/CD Pipeline' (yellow circle, 2 minutes ago, in progress), 'fix: resolve pip mirror index, add typing imports, and verify full pi...' (red X, 12 minutes ago, 4m 32s), 'Update README.md' (red X, Today at 21:19, 23s), and 'feat: complete end-to-end MLOps pipeline for AI vs Human news classif...' (red X, Today at 20:08, 21s).

aditi-govindu / mlops-final-term-project

Code Issues Pull requests Agents Actions Projects Wiki Security and quality Insights Settings

← MLOps CI/CD Pipeline

✓ **fix(docker): Move directories with .gitkeep for Docker build #10** Re-run all jobs

Summary

All jobs

- ✓ Code Quality & Linting
- ✓ Unit & Integration Tests
- ✓ E2E Pipeline Validation
- ✓ Docker Build & Smoke Test

Run details

- Usage
- Workflow file

Triggered via push 20 minutes ago

aditi-govindu pushed -> a8ea5b8 main

Status: **Success** Total duration: **14m 7s** Artifacts: **1**

ci-cd.yml
on: push

Code Quality & Linting 17% • Unit & Integration Tests 3m 19s • E2E Pipeline Validation 5m 57s • Docker Build & Smoke ... 4m 18s

Step 12: Spin up the docker image using docker-compose up -d --build

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS docker-compose
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % docker-compose up -d --build
unable to get image 'mlops-final-term-project-api': failed to connect to the docker API at unix:///Users/aditigovindu/.docker/run/docker.sock; check if the path is correct and if the daemon is running: dial unix /Users/aditigovindu/.docker/run/docker.sock: connect: no such file or directory
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % docker-compose up -d --build
WARN[0000] /Users/aditigovindu/Desktop/mlops-final-term-project/docker-compose.yml: the attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion
[+] up 41/53
.: Image apache/airflow:2.7.1-python3.9 [.....] 110.1MB / 363.6MB Pulling 215.5s
.: Image ghcr.io/mlflow/mlflow:v2.10.0 [.....] 52.51MB / 237.5MB Pulling 215.5s
.: Image prom/prometheus:v2.45.0 [.....] 41.94MB / 86.54MB Pulling 215.5s
.: Image grafana/grafana:10.0.3 [.....] 61.1MB / 86.46MB Pulling 215.5s

```

airflow-docker
 /Users/aditigovindu/kafka-itm/airflow-docker

[View configurations](#)

	airflow-worker				
	apache/airflow:3.1				
	airflow-apise...				
	apache/airflow:3.1				
	airflow-trigge...				
	apache/airflow:3.1				
	airflow-sched...				
	apache/airflow:3.1				
	airflow-dag-p...				
	apache/airflow:3.1				
	airflow-init				
	apache/airflow:3.1				
	postgres				
	postgres:16				
	redis				
	redis:7.2-bookwor				

```

dag_id=child_dag loc=collection.py:194
2026-08-19T19:51:41.408027Z [info    ] no latest
run found
[airflow.dag_processing.collection]
dag_id=child_dag loc=collection.py:204
2026-08-19T19:51:41.411376Z [info    ] get
next_dagrun_info_v2
[airflow.serialization.definitions.dag]
last_automated_run_info=None loc=dag.py:420
next_info=DagRunInfo(run_after=DateTime(2022, 1, 1,
0, 0, 0, tzinfo=Timezone('UTC')),
data_interval=DataInterval(start=DateTime(2022, 1,
1, 0, 0, 0, tzinfo=Timezone('UTC')),
end=DateTime(2022, 1, 1, 0, 0, 0,
tzinfo=Timezone('UTC'))), partition_date=None,
partition_key=None)
2026-08-19T19:51:41.411809Z [info    ] setting
next dagrun info [airflow.models.dag]
loc=dag.py:792 next_dagrun='2022-01-01
00:00:00+00:00' next_dagrun_create_after='2022-01-
01 00:00:00+00:00'
next_dagrun_data_interval_end='2022-01-01
00:00:00+00:00'
next_dagrun_data_interval_start='2022-01-01
00:00:00+00:00' n
next_dagrun_parti

```

Downloading version 4.87.0

RAM 3.36 GB CPU 10.30% Disk: 7.82 GB used (limit 223.63 GB)

Downloading

Wait for docker to pull all images

```

PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS
docker-compose

(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % docker-compose up -d --build
(venvml) aditigovindu@Aditis-MacBook-Air mlops-final-term-project % docker-compose up -d --build
WARN[0000] /Users/aditigovindu/Desktop/mlops-final-term-project/docker-compose.yml: the attribute `version`
is obsolete, it will be ignored, please remove it to avoid potential confusion
[+] up 53/53
✓ Image apache/airflow:2.7.1-python3.9 Pulled 695.3s
✓ Image ghcr.io/mlflow/mlflow:v2.10.0 Pulled 718.2s
✓ Image prom/prometheus:v2.45.0 Pulled 358.8s
✓ Image grafana/grafana:10.0.3 Pulled 361.8s
[+] Building 20.0s (5/17)
=> [internal] load metadata for docker.io/library/python:3.9-slim 5.4s
=> [internal] load .dockerignore 0.0s
=> => transferring context: 2B 0.0s

```

Step 13: Navigate to MLFlow to view the model endpoint
 mlflow ui --host 0.0.0.0 --port 5000 --backend-store-uri
 sqlite:///mlruns.db

localhost:5000/#/experiments/1/runs/685c7ec991e74e409b17f1a951d8fc33

mlflow 3.1.4 Experiments Models Prompts

ai-vs-human-news-classification >

tfidf_logistic_regression

Overview Model metrics System metrics Traces Artifacts

Description [✎](#)
No description

Details

Created at	08/18/2026, 11:13:12 PM
Created by	aditigovindu
Experiment ID	1 🔗
Status	🟢 Finished
Run ID	685c7ec991e74e409b17f1a951d8fc33 🔗
Duration	102ms
Datasets used	—
Tags	Add tags
Source	🔗 train.py
Logged models	—
Registered models	—
Registered prompts	—

View the runs on MLFlow for this experiment

localhost:5000/#/experiments/1?searchFilter=&orderByKey=attributes.start_time&orderByAsc=false&startTime=ALL&lifecycleFilter=Active&modelVersionFilter=All+Runs&datasetsFilter=...

mlflow 3.1.4 Experiments Models Prompts

ai-vs-human-news-classification [🔗](#) [Provide Feedback](#) [Add Description](#) [Share](#)

Experiments [🔍](#) [📄](#)

Search experiments

☐ Default [✎](#) [🗑️](#)

☒ ai-vs-human-news-classification [✎](#) [🗑️](#)

Runs Models Experimental Evaluation [🔍](#) Traces

[📄](#) [📊](#) [📁](#) [🔍](#) Time created [⌵](#) State: Active [⌵](#) Datasets [⌵](#) Sort: Created [⌵](#) [+ New run](#)


[📄](#) Columns [📄](#) Group by [⌵](#)

☐	Run Name	Created 📅	Dataset	Duration	Source	Models
☐	🟢 bilstm_attention_classif...	🟢 3 minutes ago	-	13.3s	🔗 train.py	-
☐	🟢 tfidf_logistic_regression	🟢 3 minutes ago	-	146ms	🔗 train.py	-
☐	🟢 bilstm_attention_classif...	🟢 6 minutes ago	-	13.4s	🔗 train.py	-
☐	🔴 tfidf_logistic_regression	🟢 6 minutes ago	-	146ms	🔗 train.py	-
☐	🟢 bilstm_attention_classif...	🟢 19 minutes ago	-	13.3s	🔗 train.py	-
☐	🟢 tfidf_logistic_regression	🟢 19 minutes ago	-	149ms	🔗 train.py	-
☐	🔴 bilstm_attention_classif...	🟢 20 minutes ago	-	13.5s	🔗 train.py	-
☐	🟢 Hidf_logistic_regression	🟢 20 minutes ago	-	148ms	🔗 train.py	-
☐	🟢 bilstm_attention_classif...	🟢 33 minutes ago	-	12.6s	🔗 train.py	-
☐	🟢 Hidf_logistic_regression	🟢 33 minutes ago	-	102ms	🔗 train.py	-

[🔍](#) Show more columns (31 total)

10 matching runs

View metrics on MLFlow



Step 14: Navigate to 0.0.0.0:5000 to view FastAPI real time serving

```
Problems  Output  Debug Console  Terminal  Ports
Python + v @ [Icons] [Close]

(venv) aditigovindu-mac:mlops-final-term-project aditigovindu$ uvicorn src.api.app:app --host 0.0.0.0 --port 8000 --reload
INFO: Will watch for changes in these directories: ['/Users/aditigovindu/Desktop/mlops-assignment/mlops-final-term-project']
INFO: Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
INFO: Started reloader process [96989] using WatchFiles
INFO: Started server process [96991]
INFO: Waiting for application startup.
[2026-08-19 00:01:43] [INFO] [api_service:68]: Initializing API application and loading model artifacts...
[2026-08-19 00:01:43] [INFO] [api_service:77]: Loading PyTorch model checkpoint from models/pytorch_bilstm.pt
[2026-08-19 00:01:43] [INFO] [api_service:98]: PyTorch model loaded successfully.
[2026-08-19 00:01:43] [INFO] [api_service:103]: Loading baseline model from models/baseline_model.pkl
[2026-08-19 00:01:43] [INFO] [api_service:113]: Drift detector initialized with reference training dataset.
INFO: Application startup complete.
INFO: 127.0.0.1:56613 - "GET / HTTP/1.1" 200 OK
INFO: 127.0.0.1:60808 - "GET /docs HTTP/1.1" 200 OK
INFO: 127.0.0.1:60808 - "GET /openapi.json HTTP/1.1" 200 OK
INFO: 127.0.0.1:49500 - "POST /predict HTTP/1.1" 200 OK
INFO: 127.0.0.1:49500 - "POST /predict HTTP/1.1" 200 OK
INFO: 127.0.0.1:64327 - "POST /predict HTTP/1.1" 200 OK
INFO: 127.0.0.1:56751 - "GET /metrics HTTP/1.1" 200 OK
INFO: 127.0.0.1:54053 - "GET /docs HTTP/1.1" 200 OK
INFO: 127.0.0.1:54053 - "GET /openapi.json HTTP/1.1" 200 OK
INFO: 127.0.0.1:62487 - "POST /predict HTTP/1.1" 200 OK
INFO: 127.0.0.1:62487 - "POST /predict HTTP/1.1" 200 OK
INFO: 127.0.0.1:57616 - "POST /predict HTTP/1.1" 200 OK
```

Step 15: Navigate to Prometheus to view model metrics at port 8000.

```
← → ↺ [ ] ⚠ Not Secure 0.0.0.0:8000/metrics

# HELP python_gc_objects_collected_total Objects collected during gc
# TYPE python_gc_objects_collected_total counter
python_gc_objects_collected_total{generation="0"} 11073.0
python_gc_objects_collected_total{generation="1"} 1499.0
python_gc_objects_collected_total{generation="2"} 39.0
# HELP python_gc_objects_uncollectable_total Uncollectable objects found during GC
# TYPE python_gc_objects_uncollectable_total counter
python_gc_objects_uncollectable_total{generation="0"} 0.0
python_gc_objects_uncollectable_total{generation="1"} 0.0
python_gc_objects_uncollectable_total{generation="2"} 0.0
# HELP python_gc_collections_total Number of times this generation was collected
# TYPE python_gc_collections_total counter
python_gc_collections_total{generation="0"} 648.0
python_gc_collections_total{generation="1"} 58.0
python_gc_collections_total{generation="2"} 5.0
# HELP python_info Python platform information
# TYPE python_info gauge
python_info{implementation="CPython",major="3",minor="9",patchlevel="6",version="3.9.6"} 1.0
# HELP http_requests_total Total HTTP Requests
# TYPE http_requests_total counter
http_requests_total{endpoint="/predict",method="POST",status_code="200"} 3.0
# HELP http_requests_created Total HTTP Requests
# TYPE http_requests_created gauge
http_requests_created{endpoint="/predict",method="POST",status_code="200"} 1.7870779448156428e+09
# HELP http_request_duration_seconds HTTP Request Latency in seconds
# TYPE http_request_duration_seconds histogram
http_request_duration_seconds_bucket{endpoint="/predict",le="0.005"} 0.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.01"} 0.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.025"} 0.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.05"} 2.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.075"} 2.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.1"} 2.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.25"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.5"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="0.75"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="1.0"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="2.5"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="5.0"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="7.5"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="10.0"} 3.0
http_request_duration_seconds_bucket{endpoint="/predict",le="+Inf"} 3.0
http_request_duration_seconds_count{endpoint="/predict"} 3.0
http_request_duration_seconds_sum{endpoint="/predict"} 0.16798049999999876
# HELP http_request_duration_seconds_created HTTP Request Latency in seconds
# TYPE http_request_duration_seconds_created gauge
http_request_duration_seconds_created{endpoint="/predict"} 1.787077944711355e+09
# HELP predictions_total Total predictions broken down by class
# TYPE predictions_total counter
predictions_total{predicted_label="Human-written"} 3.0
# HELP predictions_created Total predictions broken down by class
# TYPE predictions_created gauge
predictions_created{predicted_label="Human-written"} 1.7870779448156002e+09
# HELP prediction_confidence_score Distribution of prediction confidence scores
# TYPE prediction_confidence_score histogram
prediction_confidence_score_bucket{le="0.5"} 0.0
prediction_confidence_score_bucket{le="0.6"} 2.0
prediction_confidence_score_bucket{le="0.7"} 2.0
prediction_confidence_score_bucket{le="0.8"} 2.0
prediction_confidence_score_bucket{le="0.85"} 2.0
prediction_confidence_score_bucket{le="0.9"} 2.0
```